

文档2-Hbase2Hive

1. 基本原理

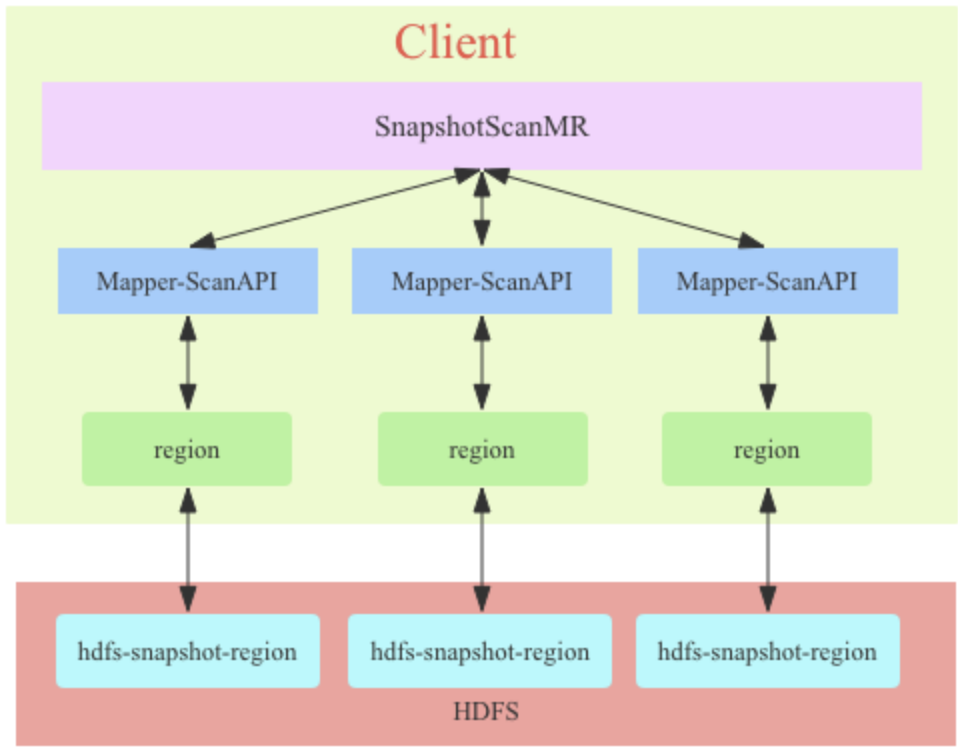
思路：

离线数据获取不影响当前Hbase业务，如果使用传统的数据扫描，将会对于Hbase RegionServer产生较大的性能压力，所以采取Snapshot方式扫描Hbase存储的数据

SnapshotScanMR直接会在客户端打开Region扫描HDFS上的文件，不需要发送Scan请求给RegionServer，再由RegionServer扫描HDFS上的文件。

优点：

- 减小对RegionServer的影响，SnapshotScanMR这种绕过RegionServer的实现方式最大限度的减小了对集群中其他业务的影响。
- 极大的提升了扫描效率。SnapshotScanMR相比TableScanMR在扫描效率上会有2倍~N倍的性能提升



应用场景：

- 数据从Hbase导出至HDFS或Hive

2. 代码示例

2.1 pom引入

Pom文件

```
<dependency>
    <groupId>org.apache.hbase</groupId>
    <artifactId>hbase-mapreduce</artifactId>
    <version>2.1.0-cdh6.3.2</version>
    <scope>compile</scope>
    <exclusions>
        <exclusion>
            <artifactId>hbase-server</artifactId>
            <groupId>org.apache.hbase</groupId>
        </exclusion>
    </exclusions>
</dependency>

<dependency>
    <groupId>org.apache.hbase</groupId>
    <artifactId>hbase-server</artifactId>
    <version>2.1.0-cdh6.3.2</version>
</dependency>

<dependency>
    <groupId>org.apache.phoenix</groupId>
    <artifactId>phoenix-core</artifactId>
    <version>5.0.0-HBase-2.0</version>
    <exclusions>
        <exclusion>
            <artifactId>hbase-server</artifactId>
            <groupId>org.apache.hbase</groupId>
        </exclusion>
    </exclusions>
</dependency>
```

2.2 代码示例

SnapshotMR

```
package com.baijiahulian.bdg;

import org.apache.commons.lang.ArrayUtils;
import org.apache.commons.lang.StringUtils;
import org.apache.commons.logging.Log;
import org.apache.commons.logging.LogFactory;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.fs.Trash;
import org.apache.hadoop.hbase.Cell;
import org.apache.hadoop.hbase.CellScanner;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.TableName;
import org.apache.hadoop.hbase.client.*;
import org.apache.hadoop.hbase.io.ImmutableBytesWritable;
import org.apache.hadoop.hbase.mapred.TableMap;
import org.apache.hadoop.hbase.mapred.TableMapReduceUtil;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.security.UserGroupInformation;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

import java.io.IOException;
import java.security.PrivilegedExceptionAction;
```

```

import java.text.ParseException;
import java.text.SimpleDateFormat;
import java.util.UUID;

/**
 * -----
 * :
 * MR-HbaseDemo-SnapShotHbase
 * <p>
 * :
 * param1:Hbase TableName
 * param2:Hive OutPutPath
 * -----
 */
public class ScanSnapShotMR extends ToolRunner implements Tool {
    // ###
    public static final Log LOG = LogFactory.getLog(ScanSnapShotMR.class);

    // ### (Hbase)
    private static final String HDFS_FS_DEFAULTFS = "fs.defaultFS";
    private static final String HDFS_FS_DEFAULTFS_VALUE = "hdfs://xxxx:8020/";

    // ### Hbase-ZKHbaseZK
    private static final String HBASE_NAME = "hbase.zookeeper.quorum";
    private static final String HBASE_VALUE = "xxx1:2181;xxx2:2181;xxx3:2181";

    // ### Hbase-
    private static final String HBASE_ROOT_DIR = "hbase.rootdir";
    private static final String HBASE_ROOT_DIR_VALUE = "hdfs://xxxx:8020/hbase";

    // ### AppMaster
    private static final String YARN_APP_MAPREDUCE_AM_STAGING_DIR = "yarn.app.mapreduce.am.staging-dir";

    // ### Hbase-io.serializations
    private static final String HDFS_IO_SERIALIZATIONS = "io.serializations";
    private static final String HDFS_IO_SERIALIZATIONS_VALUE = "org.apache.hadoop.io.serializer.
WritableSerialization,org.apache.hadoop.io.serializer.avro.AvroSpecificSerialization,org.apache.hadoop.io.
serializer.avro.AvroReflectSerialization,org.apache.hadoop.hbase.mapreduce.MutationSerialization,org.apache.
hadoop.hbase.mapreduce.ResultSerialization,org.apache.hadoop.hbase.mapreduce.KeyValueSerialization";

    // ###
    private static final String HBASE_CF = "info";

    // ### StageDir
    private static final String STAGE_DIR_PREFIX = "hdfs://xxxx:8020/tmp/hbase/stagedir/";

    // ### tmpDir
    private static final String TEP_DIR_PREFIX = "hdfs://xxxx:8020/tmp/hbase/region_tmp/";

    // ### trash stage dir
    private static final String STAGE_TRASH_DIR_PREFIX = "hdfs://xxxx:8020/user/hbase_tmp/.Trash/";

    /**
     * SnapShotMapper
     */
    static class TableSnapshotMapper extends MapReduceBase implements TableMap<Text, NullWritable> {
        // ###
        private Text mkey = new Text();

        public void map(ImmutableBytesWritable key, Result result,
            OutputCollector<Text, NullWritable> collector, Reporter reporter)
            throws IOException {
            // ###
            StringBuffer stringBuffer = new StringBuffer();
            CellScanner cellScanner = result.cellScanner();

            SimpleDateFormat format = new SimpleDateFormat("yyyyMMdd");
            boolean flag = false;
            stringBuffer.append("7");
            // rowkeyCellScanner
            while (cellScanner.advance()) {

```

```

        Cell cell = cellScanner.current();
        String qualifier = Bytes.toString(cell.getQualifierArray(), cell.getQualifierOffset(), cell.
getQualifierLength());
        String value = Bytes.toString(cell.getValueArray(), cell.getValueOffset(), cell.
getValueLength());
        try {
            if ("task_dt".equals(qualifier) && (format.parse(value).getTime() + 14 * 24 * 60 * 60 *
1000) > System.currentTimeMillis()) {
                flag = true;
            }
        } catch (ParseException e) {
            e.printStackTrace();
        }
        if ("biz_type".equals(qualifier)) {
            continue;
        }
        stringBuffer.append("\u0001" + value);
    }

    if (flag && stringBuffer.toString().split("\u0001").length == 21) {
        mkey.set(stringBuffer.toString());
        collector.collect(mkey, NullWritable.get());
    }
}

}

public int run(String[] args) throws Exception {
    int result = -1;

    Admin admin = null;
    Connection connection = null;
    String stagingDir = StringUtils.EMPTY;
    String trashStagingDir = StringUtils.EMPTY;
    FileSystem fs = null;
    Configuration hbaseConfiguration = null;

    String snapShotName = StringUtils.EMPTY;

    try {
        // -----
        // STEP 1. output_dir
        // -----

        // ###
        Configuration config = new Configuration();
        setConf(config);

        // ### output_dir
        Path path = new Path(args[1]);
        fs = path.getFileSystem(config);

        // ###
        if (fs.exists(path)) {
            LOG.info("[ScanSnapShotMR] - output_dir is exists");
        }

        // ###
        Trash trashTmp = new Trash(fs, config);
        /* if (trashTmp.moveToTrash(path)) {
            LOG.info("Moved to trash: " + path);
        } */

        // -----
        // STEP 2. HbaseSnapShot
        // -----
        hbaseConfiguration = HBaseConfiguration.create(config);
        hbaseConfiguration.set(HDFS_IO_SERIALIZATIONS, HDFS_IO_SERIALIZATIONS_VALUE);
        hbaseConfiguration.set(HBASE_ROOT_DIR, HBASE_ROOT_DIR_VALUE);
        String tableName = args[0];

```

```

//
connection = ConnectionFactory.createConnection(hbaseConfiguration);
admin = connection.getAdmin();
if (admin == null) {
    LOG.error("[ScanSnapshotMR] - Hbase Admin Error");
    return -1;
}

//
boolean exists = admin.tableExists(TableName.valueOf(tableName));
if (!exists) {
    LOG.error("[ScanSnapshotMR] - Hbase tableName = " + tableName + " is not exists.");
    return -1;
}

// snapshot
snapShotName = tableName.split(":")[1] + "_snapshot" + UUID.randomUUID().toString().replace("-",
""");

//
LOG.info("[ScanSnapshotMR] - snapShotName = " + snapShotName);

try {
    admin.snapshot(snapShotName, TableName.valueOf(tableName));
} catch (Exception ex) {
    LOG.error("[ScanSnapshotMR] - CloneSnapshot Error", ex);
    return -1;
}

// -----
// STEP 3. HbaseSnapshot
// -----
stagingDir = STAGE_DIR_PREFIX + UUID.randomUUID().toString() + "/";
trashStagingDir = STAGE_TRASH_DIR_PREFIX + stagingDir;
hbaseConfiguration.set(YARN_APP_MAPREDUCE_AM_STAGING_DIR, stagingDir);

// ### job
JobConf jobConf = new JobConf(hbaseConfiguration);

// ### Snapshot
Path tmpPath = new Path(TEP_DIR_PREFIX + tableName.replaceAll(":", "_") + UUID.randomUUID().
toString() + "/"");
Path outPutPath = new Path(args[1]);

// ### MR
jobConf.setMapperClass(TableSnapshotMapper.class);
jobConf.setMapOutputKeyClass(Text.class);
jobConf.setMapOutputValueClass(NullWritable.class);
jobConf.setJarByClass(ScanSnapshotMR.class);

// -----
// STEP 4. MR
// -----
// ###
TableMapReduceUtil.initTableSnapshotMapJob(snapShotName,
    HBASE_CF, TableSnapshotMapper.class, Text.class,
    NullWritable.class, jobConf, false, tmpPath);

// ### outputformat
FileOutputFormat.setOutputPath(jobConf, outPutPath);
TableMapReduceUtil.addDependencyJars(jobConf);

// ### job
RunningJob job = JobClient.runJob(jobConf);
job.waitForCompletion();

// ### result
result = (job.isSuccessful() == true) ? 1 : 0;

} catch (Exception ex) {
    LOG.error("[ScanSnapshotMR] - ex = {}", ex);
}

```

```

        result = -1;
    } finally {
        try {
            // ###
            Path path = new Path(stagingDir);
            fs = path.getFileSystem(hbaseConfiguration);
            Trash trashTmp = new Trash(fs, hbaseConfiguration);
            // ###
            if (fs.exists(path)) {
                /* if (trashTmp.moveToTrash(path)) {
                    LOG.info("Moved StagingDir to trash: " + path);
                }*/
            }
            // ### snapshot
            LOG.info("[ScanSnapShotMR] - SnapShotName = " + snapShotName);
            admin.deleteSnapshot(snapShotName);

            // ###
            admin.close();
            connection.close();
        } catch (Exception ex2) {
            LOG.error("[ScanSnapShotMR] - finally ex = {}", ex2);
            result = -1;
        }
    }

    // ###
    return result;
}

public void setConf(Configuration configuration) {
    // ### Hbase
    configuration.set(HBASE_NAME, HBASE_VALUE);

    // ###
    configuration.set(HDFS_FS_DEFAULTFS, HDFS_FS_DEFAULTFS_VALUE);
}

public Configuration getConf() {
    return null;
}

/**
 * [main] - Main
 * 1.1Hbase
 * 2.2
 *
 * @param args
 * @throws Exception
 */
public static void main(final String[] args) throws Exception {

    // ###
    if (ArrayUtils.isEmpty(args)) {
        LOG.error("[ScanSnapShotMR] - args is not empty!");
        System.exit(-1);
    }

    // ###
    if (args.length < 2 || args.length > 2) {
        LOG.error("[ScanSnapShotMR] - args size must be 2 ; 1) Hbase Table Name 2) Hbase Table Output Dir");
    }

    // ###
    for (int i = 0; i < args.length; i++) {
        if (LOG.isInfoEnabled()) {
            LOG.info("[ScanSnapShotMR] - arg(" + i + ")=" + args[i]);
        }
    }

    // ###

```

```
        ToolRunner.run(HBaseConfiguration.create(), new ScanSnapshotMR(), args);
    }
}
```

2.3 配置说明

2.3.1 测试环境

- 测试环境用于进行test\beta测试
- 测试环境zk
 - al-bj-bigdata-inf-hbase-test03.inf.bdg.baijiahulian:2181;al-bj-bigdata-inf-hbase-test02.inf.bdg.baijiahulian:2181;al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian:2181

域名	IP
al-bj-bigdata-inf-hbase-test03.inf.bdg.baijiahulian	172.16.18.71
al-bj-bigdata-inf-hbase-test02.inf.bdg.baijiahulian	172.16.18.69
al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian	172.16.18.68

- 测试环境HDFS
 - private static final String HDFS_FS_DEFAULTFS_VALUE = "hdfs://al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian:8020/";
 - private static final String STAGE_DIR_PREFIX = "hdfs://al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian:8020/tmp/hbase/stagedir/";
 - private static final String HBASE_ROOT_DIR_VALUE = "hdfs://al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian:8020/hbase/";
 - private static final String TEP_DIR_PREFIX = "hdfs://al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian:8020/tmp/hbase/region_tmp/";
 - private static final String STAGE_TRASH_DIR_PREFIX = "hdfs://al-bj-bigdata-inf-hbase-test01.inf.bdg.baijiahulian:8020/user/hbase_tmp/.Trash/";

2.3.2 线上环境

- 线上环境用于进行prod上线
- 具体的配置